

Convex Optimization

AJ Friend, Nick Henderson
(w/ material from Stephen Boyd and Steven Diamond)
Stanford University

June 28, 2015

Mathematical optimization

Outline

Mathematical optimization

Modeling with convexity

Image in-painting

Fault detection

Vehicle tracking

Conclusions

This talk

- ▶ high-level (biased) overview of convex optimization
- ▶ (fancy?) examples
- ▶ why *restrict* yourself to using convex optimization?

Mathematical optimization

optimization problem has form

$$\begin{array}{ll}\text{minimize} & f_0(x) \\ \text{subject to} & f_i(x) \leq 0, \quad i = 1, \dots, m\end{array}$$

- ▶ $x \in \mathbf{R}^n$ is **decision variable** (to be found)
- ▶ f_0 is objective function; f_i are constraint functions
- ▶ problem data are hid inside f_0, \dots, f_m
- ▶ variations: add equality constraints, maximize a utility function, satisfaction (feasibility), optimal trade off

Convex optimization

convex optimization problem has form


$$\begin{array}{ll}\text{minimize} & f_0(x) \\ \text{subject to} & f_i(x) \leq 0, \quad i = 1, \dots, m\end{array}$$

► f_0, \dots, f_m **convex**: for $\theta \in [0, 1]$,


$$f_i(\theta x + (1 - \theta)y) \leq \theta f_i(x) + (1 - \theta)f_i(y)$$

i.e., f_i have nonnegative (upward) curvature

Non-convex function



Convex function



Why convexity?

- ▶ local minimizers are global
- ▶ most optimization problems often cannot be solved; **convex** problems (usually) can be
- ▶ useful theory of convexity
- ▶ effective algorithms and available software
 - ▶ global solutions
 - ▶ polynomial complexity
 - ▶ algorithms that scale
- ▶ convenient **language** to discuss problems
- ▶ unifies many methods; subroutine for non-convex problems
- ▶ expressive: **lots of applications**

Applications

- ▶ machine learning, statistics
- ▶ finance
- ▶ supply chain, revenue management, advertising
- ▶ control
- ▶ signal and image processing, vision
- ▶ networking
- ▶ circuit design
- ▶ combinatorial optimization
- ▶ quantum mechanics
- ▶ flux-based analysis

Modeling with convexity

Outline

Mathematical optimization

Modeling with convexity

Image in-painting

Fault detection

Vehicle tracking

Conclusions

Approach

- ▶ try to formulate your problem as convex
- ▶ **build** model guaranteed to be convex from convex atoms and composition rules
- ▶ if you succeed, you can (usually) solve it (numerically)
- ▶ if non-convex, approximations can work surprisingly well
- ▶ an **interface**: say what you want, not how to get it

CVXPY

- ▶ we'll use CVXPY (www.cvxpy.org), convex modeling and solving tool in Python
- ▶ write code very close to the math
- ▶ CVXPY and similar tools allow for rapid prototyping
- ▶ does the work for you:
 - ▶ checks convexity
 - ▶ transforms problem (no matrix stuffing by hand!)
 - ▶ interface to several solvers

CVXPY example

math: (constrained LASSO)

$$\begin{array}{ll}\text{minimize} & \|Ax - b\|_2^2 + \rho\|x\|_1 \\ \text{subject to} & \mathbf{1}^T x = 0, \quad \|x\|_\infty \leq 1\end{array}$$

with variable $x \in \mathbf{R}^n$

code:

```
from cvxpy import *
x = Variable(n)
obj = sum_squares(A*x-b) + rho*norm(x,1)
constr = [sum_entries(x)==1, norm(x,'inf')<=1]
Problem(obj,constr).solve()
```

Summary

- ▶ good trade-off between algorithmic supervision and modeling power
- ▶ lots of theoretical, algorithmic, and software tools
- ▶ with proper training, “your problem is convex” offers deep, cosmic relief

Image in-painting

Outline

Mathematical optimization

Modeling with convexity

Image in-painting

Fault detection

Vehicle tracking

Conclusions

Image in-painting

Original



Corrupted

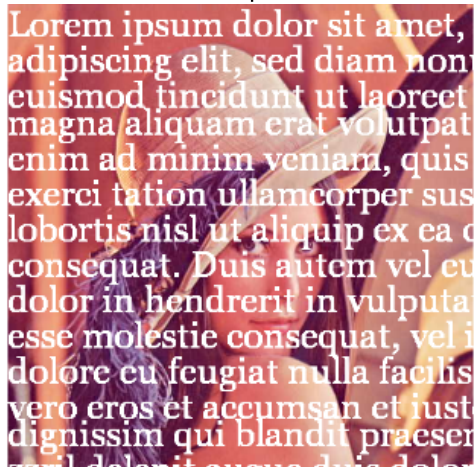


Image in-painting

guess pixel values in obscured/corrupted parts of image

- ▶ **decision variable** $x \in \mathbf{R}^{m \times n \times 3}$
- ▶ $x_{i,j} \in [0, 1]^3$ gives RGB values of pixel (i, j)
- ▶ many pixels missing
- ▶ known pixel IDs given by set K , values given by **data** $y \in \mathbf{R}^{m \times n \times 3}$

total variation in-painting: choose pixel values $x_{i,j} \in \mathbf{R}^3$ to minimize

$$\text{TV}(x) = \sum_{i,j} \left\| \begin{bmatrix} x_{i+1,j} - x_{i,j} \\ x_{i,j+1} - x_{i,j} \end{bmatrix} \right\|_2$$

that is, for each pixel, minimize distance to neighbors below and to the right, subject to known pixel values

Convex model

$$\begin{array}{ll}\text{minimize} & \text{TV}(x) \\ \text{subject to} & x_{i,j} = y_{i,j} \text{ if } (i,j) \in K\end{array}$$

- ▶ write what you want, not how to get it
- ▶ problem concisely and elegantly expressed
 - ▶ easily communicated
 - ▶ low overhead to tweaking model (rapid prototyping)
- ▶ we're done! (well, sort of)
 - ▶ express in code
 - ▶ invoke convex solver
 - ▶ made easier with model-and-solve tool, e.g., CVXPY
 - ▶ use other solvers/algorithms for speed or scale, if needed


Code example

```
# K[i, j] == 1 if pixel value known, 0 if unknown
from cvxpy import *
variables = []
constr = []
for i in range(3):
    x = Variable(rows, cols)
    variables += [x]
    constr += [mul_elemwise(K, x - y[:, :, i]) == 0]

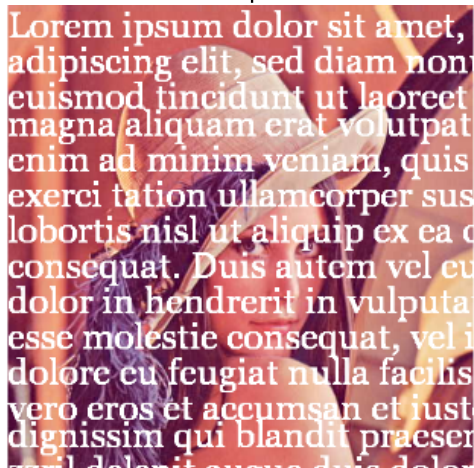
prob = Problem(Minimize(tv(*variables)), constr)
prob.solve(solver=SCS)
```

Example: 512×512 color image; about 800k variables

Original



Corrupted



Example

Original



Recovered




Example (80% of pixels removed)

Original



Corrupted




Example (80% of pixels removed)

Original



Recovered



Fault detection

Outline

Mathematical optimization

Modeling with convexity

Image in-painting

Fault detection

Vehicle tracking

Conclusions

Statistical model fitting

- ▶ fitting parameters in statistical models is often convex
 - ▶ least-squares
 - ▶ general maximum-likelihood estimation
- ▶ when not convex, approximations can work surprisingly well

Fault detection

- ▶ each of n possible faults occurs independently with probability p
- ▶ encode as $x_i \in \{0, 1\}$
- ▶ m sensors measure system performance
- ▶ sensor output is $y = Ax + v = \sum_{i=1}^n x_i a_i + v$
- ▶ v is Gaussian noise with variance σ^2
- ▶ $a_i \in \mathbf{R}^m$ is *fault signature* for fault i
- ▶ signal-to-noise ratio

$$\text{SNR} = \frac{\mathbf{E} \|Ax\|^2}{\mathbf{E} \|v\|^2}$$

- ▶ **goal:** guess x (which faults have occurred) given y (sensor measurements)

Maximum likelihood estimation

- ▶ choose $x \in \{0, 1\}^n$ to minimize negative log likelihood function

$$\ell(x) = \frac{1}{2\sigma^2} \|Ax - y\|_2^2 + \log(1/p - 1) \mathbf{1}^T x + c,$$

- ▶ nonconvex, NP-hard
- ▶ instead solve convex (relaxed) problem

$$\begin{array}{ll} \text{minimize} & \|Ax - y\|_2^2 + 2\sigma^2 \log(1/p - 1) \mathbf{1}^T x \\ \text{subject to} & 0 \leq x_i \leq 1, \quad i = 1, \dots, n \end{array}$$

and round

- ▶ called **relaxed ML** estimate

Relaxed ML CVXPY code

$$\begin{array}{ll}\text{minimize} & \|Ax - y\|_2^2 + 2\sigma^2 \log(1/p - 1) \mathbf{1}^T x \\ \text{subject to} & 0 \leq x_i \leq 1, \quad i = 1, \dots, n\end{array}$$


```
from cvxpy import *
x = Variable(n)
tau = 2*log(1/p - 1)*sigma**2
obj = Minimize(sum_squares(A*x-y) + tau*sum_entries(x))
constr = [0<=x, x<=1]
Problem(obj,constr).solve()

x_rml = np.array(x.value).flatten()
x_rnd = (x_rml>=.5).astype(int)
```


Example

$n = 2000$ possible faults, $m = 200$ measurements

$p = 0.01$, $\text{SNR} = 5$



► perfect fault recovery!

Vehicle tracking

Outline

Mathematical optimization

Modeling with convexity

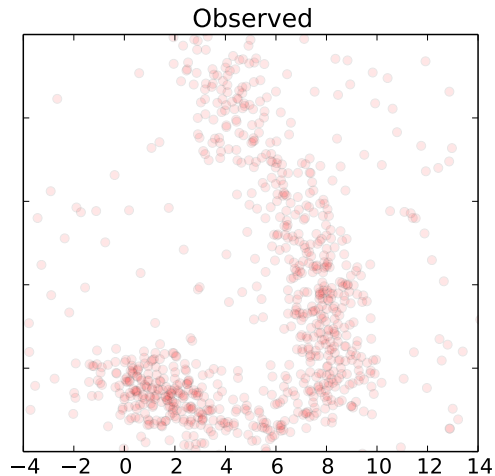

Image in-painting

Fault detection

Vehicle tracking

Conclusions

Vehicle tracking



Kalman filtering

- ▶ estimate vehicle path from noisy position measurements (with outliers)
- ▶ dynamic model of vehicle state x_t :

$$x_{t+1} = Ax_t + Bw_t, \quad y_t = Cx_t + v_t$$

- ▶ x_t is vehicle state (position, velocity)
 - ▶ w_t is unknown drive force on vehicle
 - ▶ y_t is position measurement; v_t is noise
- ▶ Kalman filter: estimate x_t by minimizing $\sum_t (\|w_t\|_2^2 + \gamma\|v_t\|_2^2)$
- ▶ a least-squares problem; assumes w_t, v_t Gaussian

Robust Kalman filter

- ▶ to handle outliers in v_t , replace square cost with Huber cost
- ▶ **robust** Kalman filter:


$$\begin{aligned} & \text{minimize} && \sum_t (\|w_t\|_2^2 + \gamma\phi(v_t)) \\ & \text{subject to} && x_{t+1} = Ax_t + Bw_t, \quad y_t = Cx_t + v_t \end{aligned}$$

where ϕ is Huber function

$$\phi(a) = \begin{cases} \|a\|_2^2 & \|a\|_2 \leq 1 \\ 2\|a\| - 1 & \|a\|_2 > 1 \end{cases}$$

- ▶ a convex problem

Huber loss function



- ▶ outside $[-1, 1]$ interval, penalizes linearly, not quadratically
- ▶ large errors more easily “forgiven”, allowing fit to discount outliers

Robust KF CVXPY code

```
from cvxpy import *
x = Variable(4,n+1)
w = Variable(2,n)
v = Variable(2,n)

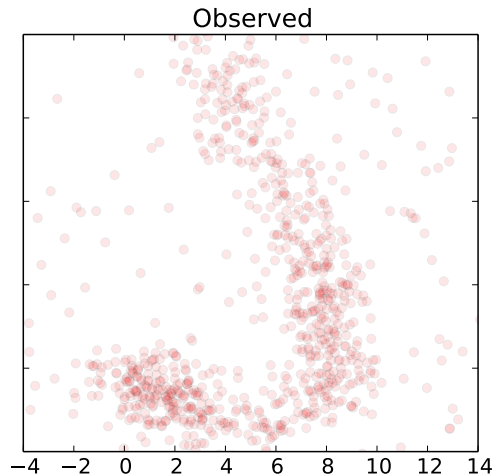
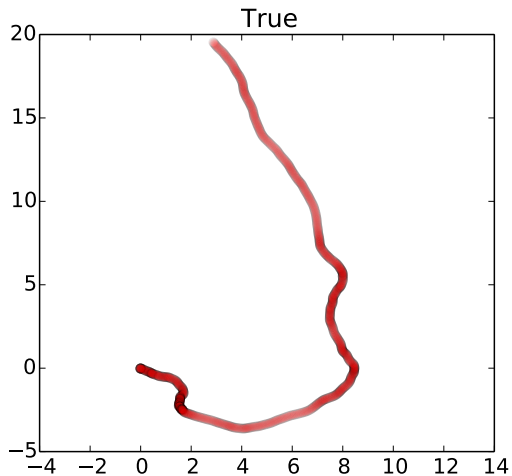
obj = sum_squares(w)
obj += sum(huber(norm(v[:,t]))) for t in range(n))
obj = Minimize(obj)
constr = []
for t in range(n):
    constr += [ x[:,t+1] == A*x[:,t] + B*w[:,t] ,
                y[:,t]   == C*x[:,t] + v[:,t]   ]

Problem(obj, constr).solve()
```




Example

- ▶ 1000 time steps
- ▶ w_t standard Gaussian
- ▶ v_t standard Gaussian, except 30% are outliers with $\sigma = 10$



Example



Example



Example



Conclusions

Outline

Mathematical optimization

Modeling with convexity

Image in-painting

Fault detection

Vehicle tracking

Conclusions

Conclusions

- ▶ convex optimization problems arise in many applications
- ▶ can be solved effectively
- ▶ high level languages (CVX, CVXPY) make prototyping easy

but lots we couldn't cover in this presentation:

- ▶ theory of convexity
- ▶ duality/Lagrange multipliers
- ▶ algorithms
 - ▶ (stochastic) gradient descent
 - ▶ interior point algorithms
 - ▶ optimal first order algorithms
 - ▶ problem splitting for large scale and distributed optimization

References

- ▶ *Convex Optimization* (Boyd & Vandenberghe)
- ▶ *CVX: Matlab software for disciplined convex programming* (Grant & Boyd)
- ▶ CVXPY (www.cvxpy.org)