

Convex Optimization

Stephen Boyd

Electrical Engineering
Computer Science
Management Science and Engineering
Institute for Computational Mathematics & Engineering
Stanford University

MLSS, Tübingen, July 15-17 2015

Outline

Mathematical Optimization

Convex Optimization

Examples

Real-Time Embedded Optimization

Large-Scale Distributed Optimization

Summary

Outline

Mathematical Optimization

Convex Optimization

Examples

Real-Time Embedded Optimization

Large-Scale Distributed Optimization

Summary

Optimization problem

$$\begin{array}{ll}\text{minimize} & f_0(x) \\ \text{subject to} & f_i(x) \leq 0, \quad i = 1, \dots, m \\ & g_i(x) = 0, \quad i = 1, \dots, p\end{array}$$

- ▶ $x \in \mathbf{R}^n$ is (vector) variable to be chosen
- ▶ f_0 is the *objective function*, to be minimized
- ▶ f_1, \dots, f_m are the *inequality constraint functions*
- ▶ g_1, \dots, g_p are the *equality constraint functions*
- ▶ variations: maximize objective, multiple objectives, ...

Finding good (or best) actions

- ▶ x represents some *action*, e.g.,
 - ▶ trades in a portfolio
 - ▶ airplane control surface deflections
 - ▶ schedule or assignment
 - ▶ resource allocation
 - ▶ transmitted signal
- ▶ constraints limit actions or impose conditions on outcome
- ▶ the smaller the objective $f_0(x)$, the better
 - ▶ total cost (or negative profit)
 - ▶ deviation from desired or target outcome
 - ▶ risk
 - ▶ fuel use

Engineering design

- ▶ x represents a design (of a circuit, device, structure, ...)
- ▶ constraints come from
 - ▶ manufacturing process
 - ▶ performance requirements
- ▶ objective $f_0(x)$ is combination of cost, weight, power, ...

Finding good models

- ▶ x represents the *parameters* in a model
- ▶ constraints impose requirements on model parameters (e.g., nonnegativity)
- ▶ objective $f_0(x)$ is the prediction error on some observed data (and possibly a term that penalizes model complexity)

Inversion

- ▶ x is something we want to estimate/reconstruct, given some measurement y
- ▶ constraints come from prior knowledge about x
- ▶ objective $f_0(x)$ measures deviation between predicted and actual measurements

Worst-case analysis (pessimization)

- ▶ variables are actions or parameters out of our control (and possibly under the control of an adversary)
- ▶ constraints limit the possible values of the parameters
- ▶ minimizing $-f_0(x)$ finds *worst possible parameter values*
- ▶ if the worst possible value of $f_0(x)$ is tolerable, you're OK
- ▶ it's good to know what the worst possible scenario can be

Optimization-based models

- ▶ model an entity as taking actions that solve an optimization problem
 - ▶ an individual makes choices that maximize expected utility
 - ▶ an organism acts to maximize its reproductive success
 - ▶ reaction rates in a cell maximize growth
 - ▶ currents in a circuit minimize total power

Optimization-based models

- ▶ model an entity as taking actions that solve an optimization problem
 - ▶ an individual makes choices that maximize expected utility
 - ▶ an organism acts to maximize its reproductive success
 - ▶ reaction rates in a cell maximize growth
 - ▶ currents in a circuit minimize total power
- ▶ (except the last) these are *very crude* models
- ▶ and yet, they often work very well

Summary

- ▶ **summary:** optimization arises *everywhere*

Summary

- ▶ **summary:** optimization arises *everywhere*
- ▶ **the bad news:** most optimization problems are *intractable*
i.e., we cannot solve them

Summary

- ▶ **summary:** optimization arises *everywhere*
- ▶ **the bad news:** most optimization problems are *intractable*
i.e., we cannot solve them
- ▶ **an exception:** *convex optimization problems are tractable*
i.e., we (generally) can solve them

Outline

Mathematical Optimization

Convex Optimization

Examples

Real-Time Embedded Optimization

Large-Scale Distributed Optimization

Summary

Convex optimization — Classical form

convex optimization problem:

$$\begin{array}{ll}\text{minimize} & f_0(x) \\ \text{subject to} & f_i(x) \leq 0, \quad i = 1, \dots, m \\ & Ax = b\end{array}$$

- ▶ variable $x \in \mathbf{R}^n$
- ▶ equality constraints are linear
- ▶ f_0, \dots, f_m are **convex**: for $\theta \in [0, 1]$,

$$f_i(\theta x + (1 - \theta)y) \leq \theta f_i(x) + (1 - \theta)f_i(y)$$

i.e., f_i have nonnegative (upward) curvature

Convex optimization — Cone form

$$\begin{array}{ll}\text{minimize} & c^T x \\ \text{subject to} & x \in K \\ & Ax = b\end{array}$$

- ▶ variable $x \in \mathbf{R}^n$
- ▶ $K \subset \mathbf{R}^n$ is a proper cone
 - ▶ K nonnegative orthant \rightarrow LP
 - ▶ K Lorentz cone \rightarrow SOCP
 - ▶ K positive semidefinite matrices \rightarrow SDP
- ▶ the ‘modern’ canonical form

Why

- ▶ beautiful, nearly complete theory
 - ▶ duality, optimality conditions, ...

Why

- ▶ beautiful, nearly complete theory
 - ▶ duality, optimality conditions, ...
- ▶ effective algorithms, methods (in theory and practice)
 - ▶ get **global solution** (and optimality certificate)
 - ▶ polynomial complexity

Why

- ▶ beautiful, nearly complete theory
 - ▶ duality, optimality conditions, ...
- ▶ effective algorithms, methods (in theory and practice)
 - ▶ get **global solution** (and optimality certificate)
 - ▶ polynomial complexity
- ▶ conceptual unification of many methods

Why

- ▶ beautiful, nearly complete theory
 - ▶ duality, optimality conditions, ...
- ▶ effective algorithms, methods (in theory and practice)
 - ▶ get **global solution** (and optimality certificate)
 - ▶ polynomial complexity
- ▶ conceptual unification of many methods
- ▶ **lots of applications** (many more than previously thought)

Application areas

- ▶ machine learning, statistics
- ▶ finance
- ▶ supply chain, revenue management, advertising
- ▶ control
- ▶ signal and image processing, vision
- ▶ networking
- ▶ circuit design
- ▶ combinatorial optimization
- ▶ quantum mechanics
- ▶ flux-based analysis

The approach

- ▶ try to formulate your optimization problem as convex
- ▶ if you succeed, you can (usually) solve it (numerically)

The approach

- ▶ try to formulate your optimization problem as convex
- ▶ if you succeed, you can (usually) solve it (numerically)
- ▶ some tricks:
 - ▶ change of variables
 - ▶ approximation of true objective, constraints
 - ▶ *relaxation*: ignore terms or constraints you can't handle

Outline

Mathematical Optimization

Convex Optimization

Examples

Real-Time Embedded Optimization

Large-Scale Distributed Optimization

Summary

Radiation treatment planning

- ▶ radiation beams with intensities x_j are directed at patient
- ▶ radiation dose y_i received in voxel i
- ▶ $y = Ax$
- ▶ $A \in \mathbf{R}^{m \times n}$ comes from beam geometry, physics
- ▶ goal is to choose x to deliver prescribed radiation dose d_i
 - ▶ $d_i = 0$ for non-tumor voxels
 - ▶ $d_i > 0$ for tumor voxels
- ▶ $y = d$ not possible, so we'll need to compromise
- ▶ typical problem has $n = 10^3$ beams, $m = 10^6$ voxels

Radiation treatment planning via convex optimization


$$\begin{array}{ll}\text{minimize} & \sum_i f_i(y_i) \\ \text{subject to} & x \geq 0, \quad y = Ax\end{array}$$

- ▶ variables $x \in \mathbf{R}^n$, $y \in \mathbf{R}^m$
- ▶ objective terms are

$$f_i(y_i) = w_i^{\text{over}}(y_i - d_i)_+ + w_i^{\text{under}}(d_i - y_i)_+$$


- ▶ w_i^{over} and w_i^{under} are positive weights
- ▶ *i.e.*, we charge linearly for over- and under-dosing
- ▶ a convex optimization problem

Example



- ▶ real patient case with $n = 360$ beams, $m = 360000$ voxels
- ▶ optimization-based plan essentially the same as plan used

Example



- ▶ real patient case with $n = 360$ beams, $m = 360000$ voxels
- ▶ optimization-based plan essentially the same as plan used
- ▶ (but we computed the plan in a few seconds, not many hours)

Image in-painting

- ▶ guess pixel values in obscured/corrupted parts of image
- ▶ *total variation in-painting*: choose pixel values $x_{ij} \in \mathbf{R}^3$ to minimize *total variation*


$$\text{TV}(x) = \sum_{ij} \left\| \begin{bmatrix} x_{i+1,j} - x_{ij} \\ x_{i,j+1} - x_{ij} \end{bmatrix} \right\|_2$$

- ▶ a convex problem

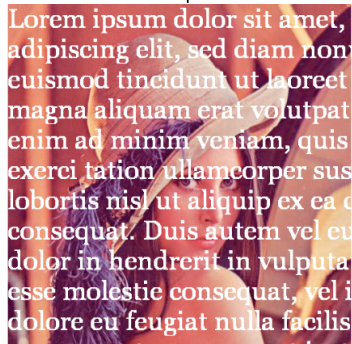
Example

512×512 color image ($n \approx 800000$ variables)

Original




Corrupted




Example

Original



Recovered



Support vector machine

- ▶ goal: predict a Boolean outcome from a set of n features
 - ▶ e.g., spam filter, fraud detection, customer purchase

Support vector machine

- ▶ goal: predict a Boolean outcome from a set of n features
 - ▶ e.g., spam filter, fraud detection, customer purchase
- ▶ data (a_i, b_i) , $i = 1, \dots, m$
 - ▶ $a_i \in \mathbf{R}^n$ feature vectors; $b_i \in \{-1, 1\}$ Boolean outcomes
- ▶ linear predictor: $\hat{b} = \text{sign}(w^T a - v)$
 - ▶ $w \in \mathbf{R}^n$ is weight vector; $v \in \mathbf{R}$ is offset

Support vector machine


- ▶ goal: predict a Boolean outcome from a set of n features
 - ▶ e.g., spam filter, fraud detection, customer purchase
- ▶ data (a_i, b_i) , $i = 1, \dots, m$
 - ▶ $a_i \in \mathbf{R}^n$ feature vectors; $b_i \in \{-1, 1\}$ Boolean outcomes
- ▶ linear predictor: $\hat{b} = \text{sign}(w^T a - v)$
 - ▶ $w \in \mathbf{R}^n$ is weight vector; $v \in \mathbf{R}$ is offset
- ▶ SVM: choose w, v to minimize (convex) objective

$$(1/m) \sum_{i=1}^m \left(1 - b_i(w^T a_i - v)\right)_+ + (\lambda/2) \|w\|_2^2$$

where $\lambda > 0$ is parameter

SVM

$$w^T z - v = 0 \text{ (solid); } |w^T z - v| = 1 \text{ (dashed)}$$



Sparsity via ℓ_1 regularization

- ▶ adding ℓ_1 -norm regularization

$$\lambda \|x\|_1 = \lambda(|x_1| + |x_2| + \cdots + |x_n|)$$

to objective results in **sparse** x

- ▶ $\lambda > 0$ controls trade-off of sparsity versus main objective
- ▶ **preserves convexity, hence tractability**
- ▶ used for many years, in many fields
 - ▶ sparse design
 - ▶ feature selection in machine learning (lasso, SVM, ...)
 - ▶ total variation reconstruction in signal processing
 - ▶ compressed sensing

Lasso

- ▶ regression problem with ℓ_1 regularization:

$$\text{minimize } (1/2)\|Ax - b\|_2^2 + \lambda\|x\|_1$$

with $A \in \mathbf{R}^{m \times n}$

- ▶ useful even when $n \gg m$ (!!); does **feature selection**

Lasso

- ▶ regression problem with ℓ_1 regularization:

$$\text{minimize} \quad (1/2)\|Ax - b\|_2^2 + \lambda\|x\|_1$$

with $A \in \mathbf{R}^{m \times n}$

- ▶ useful even when $n \gg m$ (!!); does **feature selection**
- ▶ cf. ℓ_2 regularization ('ridge regression'):

$$\text{minimize} \quad (1/2)\|Ax - b\|_2^2 + \lambda\|x\|_2^2$$

Lasso

- ▶ regression problem with ℓ_1 regularization:

$$\text{minimize } (1/2)\|Ax - b\|_2^2 + \lambda\|x\|_1$$

with $A \in \mathbf{R}^{m \times n}$

- ▶ useful even when $n \gg m$ (!!); does **feature selection**
- ▶ cf. ℓ_2 regularization ('ridge regression'):


$$\text{minimize } (1/2)\|Ax - b\|_2^2 + \lambda\|x\|_2^2$$

- ▶ lasso, ridge regression have **same computational cost**


Example

- ▶ $m = 200$ examples, $n = 1000$ features
- ▶ examples are noisy linear measurements of true x
- ▶ true x is sparse (30 nonzeros)

true x




ℓ_2 reconstruction




Example

true x



ℓ_1 (lasso) reconstruction



State of the art — Medium scale solvers

- ▶ 1000s–10000s variables, constraints
- ▶ reliably solved by interior-point methods on single machine
- ▶ exploit problem sparsity
- ▶ not quite a technology, but getting there

State of the art — Modeling languages

- ▶ (new) high level language support for convex optimization
 - ▶ describe problem in high level language
 - ▶ description is automatically transformed to cone problem
 - ▶ solved by standard solver, transformed back to original form

State of the art — Modeling languages

- ▶ (new) high level language support for convex optimization
 - ▶ describe problem in high level language
 - ▶ description is automatically transformed to cone problem
 - ▶ solved by standard solver, transformed back to original form
- ▶ enables rapid prototyping (for small and medium problems)
- ▶ ideal for teaching (can do a lot with short scripts)

CVX

- ▶ parser/solver written in Matlab (M. Grant, 2005)
- ▶ SVM: minimize

$$(1/m) \sum_{i=1}^m \left(1 - b_i(w^T a_i - v) \right)_+ + (\lambda/2) \|w\|_2^2$$

- ▶ CVX specification:

```
cvx_begin
    variables w(n) v    % weight, offset
    L=(1/m)*sum(pos(1-b.*(A*w-v)));    % avg. loss
    minimize (L+(lambda/2)*sum_square(w))
cvx_end
```

Outline

Mathematical Optimization

Convex Optimization

Examples

Real-Time Embedded Optimization

Large-Scale Distributed Optimization

Summary


Real-time embedded optimization

- ▶ in many applications, need to solve the same problem repeatedly with different data
 - ▶ control: update actions as sensor signals, goals change
 - ▶ finance: rebalance portfolio as prices, predictions change
- ▶ requires extreme solver reliability, hard real-time execution
- ▶ used now when solve times are measured in minutes, hours
 - ▶ supply chain, chemical process control, trading

Real-time embedded optimization



- ▶ in many applications, need to solve the same problem repeatedly with different data
 - ▶ control: update actions as sensor signals, goals change
 - ▶ finance: rebalance portfolio as prices, predictions change
- ▶ requires extreme solver reliability, hard real-time execution
- ▶ used now when solve times are measured in minutes, hours
 - ▶ supply chain, chemical process control, trading
- ▶ (using new techniques) can be used for applications with solve times measured in **milliseconds** or **microseconds**

Example — Positioning



- ▶ force $F(t)$ moves object, modeled as 3 masses (2 vibration modes)
- ▶ goal: move object to commanded position as quickly as possible, with $|F(t)| \leq 1$
- ▶ reduces to a (quasi-) convex problem

Optimal force profile



CVXGEN code generator

- ▶ handles small, medium size problems transformable to QP (J. Mattingley, 2010)
- ▶ accepts high-level problem family description
- ▶ uses primal-dual interior-point method
- ▶ generates flat library-free C source

CVXGEN code generator

- ▶ handles small, medium size problems transformable to QP (J. Mattingley, 2010)
- ▶ accepts high-level problem family description
- ▶ uses primal-dual interior-point method
- ▶ generates flat library-free C source

- ▶ typical speed-up over general solver: **100–10000×**

CVXGEN example specification — SVM

```
dimensions
  m = 50    % training examples
  n = 10    % dimensions
end
parameters
  a[i] (n), i = 1..m    % features
  b[i], i = 1..m    % outcomes
  lambda positive
end
variables
  w (n)    % weights
  v        % offset
end
minimize
  (1/m)*sum[i = 1..m](pos(1 - b[i]*(w'*a[i] - v))) +
  (lambda/2)*quad(w)
end
```

CVXGEN sample solve times

| problem | SVM | Positioning |
|------------------|-------------|-------------|
| variables | 61 | 590 |
| constraints | 100 | 742 |
| CVX, Intel i3 | 270 ms | 2100 ms |
| CVXGEN, Intel i3 | 230 μ s | 4.8 ms |

Outline

Mathematical Optimization

Convex Optimization

Examples

Real-Time Embedded Optimization

Large-Scale Distributed Optimization

Summary

Large-scale distributed optimization

- ▶ *large-scale* optimization problems arise in many applications
 - ▶ machine learning/statistics with huge datasets
 - ▶ dynamic optimization on large-scale networks
 - ▶ image, video processing

Large-scale distributed optimization

- ▶ *large-scale* optimization problems arise in many applications
 - ▶ machine learning/statistics with huge datasets
 - ▶ dynamic optimization on large-scale networks
 - ▶ image, video processing
- ▶ we'll use *distributed optimization*
 - ▶ split variables/constraints/objective terms among a set of *agents/processors/devices*
 - ▶ agents coordinate to solve large problem, by passing relatively small messages
 - ▶ can target modern large-scale computing platforms
 - ▶ long history, going back to 1950s

Consensus optimization

- ▶ want to solve problem with N objective terms

$$\text{minimize } \sum_{i=1}^N f_i(x)$$

e.g., f_i is the loss function for i th block of training data

- ▶ consensus form:

$$\begin{array}{ll} \text{minimize} & \sum_{i=1}^N f_i(x_i) \\ \text{subject to} & x_i - z = 0 \end{array}$$

- ▶ x_i are **local variables**
- ▶ z is the **global variable**
- ▶ $x_i - z = 0$ are **consistency** or **consensus** constraints

Consensus optimization via ADMM

with $\bar{x}^k = (1/N) \sum_{i=1}^N x_i^k$ (average over local variables)

$$x_i^{k+1} := \underset{x_i}{\operatorname{argmin}} \left(f_i(x_i) + (\rho/2) \|x_i - \bar{x}^k + u_i^k\|_2^2 \right)$$


$$u_i^{k+1} := u_i^k + (x_i^{k+1} - \bar{x}^{k+1})$$

- ▶ get **global** minimum, under very general conditions
- ▶ u^k is running sum of inconsistencies (PI control)
- ▶ minimizations carried out independently and in parallel
- ▶ coordination is via averaging of local variables x_i


Example — Consensus SVM

- ▶ baby problem with $n = 2$, $m = 400$ to illustrate
- ▶ examples split into $N = 20$ groups, in worst possible way: each group contains only positive or negative examples


Iteration 1



Iteration 5



Iteration 40



Outline

Mathematical Optimization

Convex Optimization

Examples

Real-Time Embedded Optimization

Large-Scale Distributed Optimization

Summary

Summary

- ▶ convex optimization problems **arise in many applications**
- ▶ convex optimization problems **can be solved effectively**
 - ▶ small problems at microsecond/millisecond time scales
 - ▶ medium-scale problems using general purpose methods
 - ▶ arbitrary-scale problems using distributed optimization
- ▶ high level language support (CVX) makes prototyping easy

References

many researchers have worked on the topics covered

- ▶ *Convex Optimization* (Boyd & Vandenberghe)
- ▶ *CVX: Matlab software for disciplined convex programming* (Grant & Boyd)
- ▶ *CVXGEN: A code generator for embedded convex optimization* (Mattingley & Boyd)
- ▶ *Distributed optimization and statistical learning via the alternating direction method of multipliers* (Boyd, Parikh, Chu, Peleato, & Eckstein)

all available (with code) from stanford.edu/~boyd